

## LAB 4: Exploratory Data Analysis, Feature Engineering, and Feature Selection

### Dataset

```
In [1]:
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]:
df = sns.load_dataset('titanic')
```

### Exercise1: Exploratory Data Analysis (EDA)

#### Step 1: Basic Understanding

- Use `df.info()` and `df.describe()` to understand the structure and summary
- Display column types and count of missing values

```
In [3]:
print(df.info())
print(df.describe)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   survived        891 non-null   int64  
1   pclass          891 non-null   int64  
2   sex             891 non-null   object  
3   age            714 non-null   float64 
4   sibsp          891 non-null   int64  
5   parch          891 non-null   int64  
6   fare           891 non-null   float64 
7   embarked       889 non-null   object  
8   class          891 non-null   category
9   who            891 non-null   object  
10  adult_male     891 non-null   bool    
11  deck          203 non-null   category
12  embark_town    889 non-null   object  
13  alive          891 non-null   object  
14  alone          891 non-null   bool    
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
None
<bound method NDFrame.describe of      survived  pclass    sex  age  sibsp
parch    fare embarked  class \
0         0      3   male  22.0      1      0   7.2500      S  Third
1         1      1  female  38.0      1      0  71.2833      C  First
2         1      3  female  26.0      0      0   7.9250      S  Third
```

```

3      1      1  female  35.0      1      0  53.1000      S  First
4      0      3   male  35.0      0      0   8.0500      S  Third
..      ...      ...      ...      ...      ...      ...      ...      ...
886     0      2   male  27.0      0      0  13.0000      S  Second
887     1      1  female  19.0      0      0  30.0000      S  First
888     0      3  female   NaN      1      2  23.4500      S  Third
889     1      1   male  26.0      0      0  30.0000      C  First
890     0      3   male  32.0      0      0   7.7500      Q  Third

```

```

      who  adult_male  deck  embark_town  alive  alone
0      man          True  NaN  Southampton    no  False
1  woman          False    C    Cherbourg   yes  False
2  woman          False  NaN  Southampton   yes   True
3  woman          False    C    Southampton   yes  False
4      man          True  NaN  Southampton    no   True
..      ...      ...      ...      ...      ...      ...
886   man          True  NaN  Southampton    no   True
887 woman          False    B    Southampton   yes   True
888 woman          False  NaN  Southampton    no  False
889   man          True    C    Cherbourg   yes   True
890   man          True  NaN  Queenstown    no   True

```

```
[891 rows x 15 columns]>
```

```
In [4]:
print(df.isnull().sum())
```

```

survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked       2
class         0
who           0
adult_male     0
deck         688
embark_town     2
alive          0
alone          0
dtype: int64

```

## Step 2: Identify Attribute Types

- Manually classify the following attributes:
- Categorical
- Numerical
- Target

## Step 3: Understand Distribution of Attributes

- Compute: mean, median, std, quartiles for age, fare, parch
- Plot:

- Histogram and Boxplot for age, fare
- Countplot for sex, embarked, class

```
In [5]:  
# Compute: mean, median, std, quartiles for age, fare, parch  
for x in ['age', 'fare', 'parch']:  
    print(x)  
    print("Mean: ",df[x].mean())  
    print("Median: ",df[x].median())  
    print("Std: ",df[x].std())  
    print("Quantile Q1: ",df[x].quantile(0.25))  
    print("Quantile Q2: ",df[x].quantile(0.5))  
    print("Quantile Q3: ",df[x].quantile(0.75))  
    print ("\n")
```

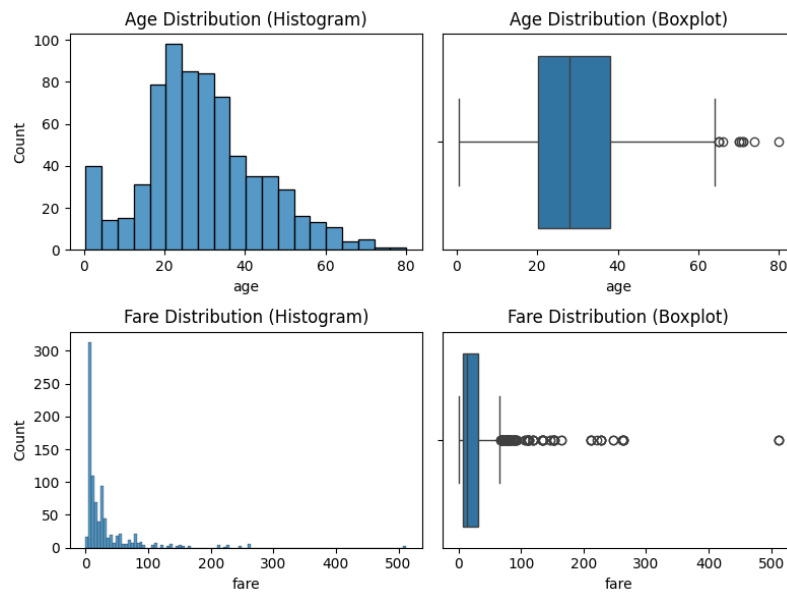
```
age  
Mean:  29.69911764705882  
Median:  28.0  
Std:  14.526497332334044  
Quantile Q1:  20.125  
Quantile Q2:  28.0  
Quantile Q3:  38.0
```

```
fare  
Mean:  32.204207968574636  
Median:  14.4542  
Std:  49.693428597180905  
Quantile Q1:  7.9104  
Quantile Q2:  14.4542  
Quantile Q3:  31.0
```

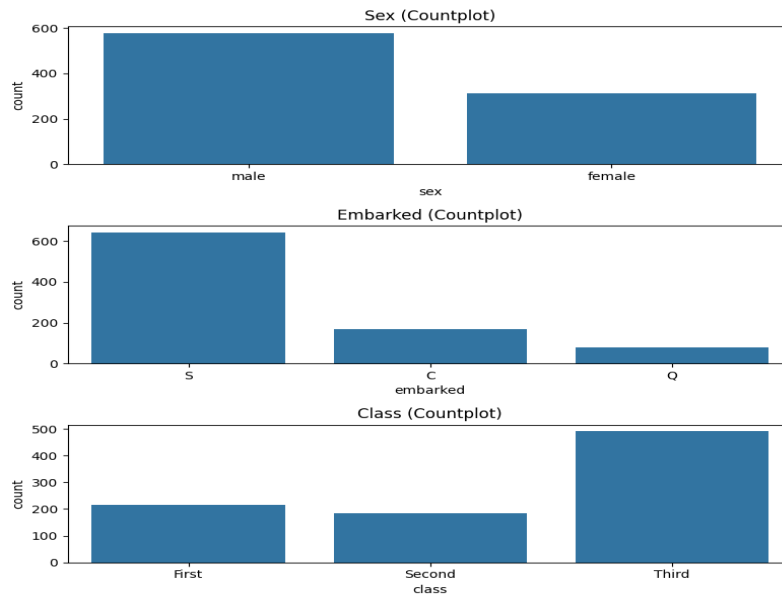
```
parch  
Mean:  0.38159371492704824  
Median:  0.0  
Std:  0.8060572211299559  
Quantile Q1:  0.0  
Quantile Q2:  0.0  
Quantile Q3:  0.0
```

```
In [6]:  
#Histogram and Boxplot for age, fare  
fig, axes = plt.subplots(2, 2, figsize=(8, 6))  
  
sns.histplot(df['age'], ax=axes[0, 0])  
axes[0, 0].set_title('Age Distribution (Histogram)')  
  
sns.boxplot(x=df['age'], ax=axes[0, 1])  
axes[0, 1].set_title('Age Distribution (Boxplot)')  
  
sns.histplot(df['fare'], ax=axes[1, 0])  
axes[1, 0].set_title('Fare Distribution (Histogram)')  
  
sns.boxplot(x=df['fare'], ax=axes[1, 1])  
axes[1, 1].set_title('Fare Distribution (Boxplot)')
```

```
plt.tight_layout()  
plt.show()
```




```
In [7]:  
# Countplot for sex, embarked, class  
  
fig, axes = plt.subplots(3, 1, figsize=(8, 8))  
  
sns.countplot(x='sex', data=df, ax=axes[0])  
axes[0].set_title('Sex (Countplot)')  
  
sns.countplot(x='embarked', data=df, ax=axes[1])  
axes[1].set_title('Embarked (Countplot)')  
  
sns.countplot(x='class', data=df, ax=axes[2])  
axes[2].set_title('Class (Countplot)')  
  
plt.tight_layout()  
plt.show()
```



#### Step 4: Understand Relationships Among Attributes

- Plot scatterplot between age and fare
- Compute and visualize Pearson correlation matrix for age, fare, parch, sibsp
- Use `pd.crosstab()` between:
  - sex vs survived
  - embarked vs class

```
In [8]:  
# Plot scatterplot between age and fare  
sns.scatterplot(x=df['age'], y = df["fare"])  
plt.show()
```



```
In [9]:  
# Compute and visualize Pearson correlation matrix for age, fare, parch, sibsp
```

```
correlation = df[['age', 'fare', 'parch', 'sibsp']].corr()
sns.heatmap(correlation, annot=True)
plt.show()
```



```
In [10]:
#Use pd.crosstab() between:
# - sex vs survived
print("sex vs survived")
pd.crosstab(df['sex'], df['survived'])
```

sex vs survived

survived	0	1
sex		
female	81	233
male	468	109

```
In [11]:
# - embarked vs class
print("embarked vs class")
pd.crosstab(df['embarked'], df['class'])
```

embarked vs class

class	First	Second	Third
embarked			
C	85	17	66
Q	2	3	72
S	127	164	353

## Exercise2: Feature Engineering

### Step 1: Missing Value Handling

- Impute:

- age with median
- embarked with mode
- Drop rows where deck is null

```
In [12]:
# age with median
df['age'].fillna(df['age'].median(), inplace=True)
# print(df['age'])
# print(df['age'].median())
```

/tmp/ipython-input-3938190663.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['age'].fillna(df['age'].median(), inplace=True)
```

```
In [13]:
# embarked with mode
df['embarked'].fillna(df['embarked'].mode()[0], inplace=True)
# print(df['embarked'])
# print(df['embarked'].mode()[0])
```

/tmp/ipython-input-3595710167.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['embarked'].fillna(df['embarked'].mode()[0], inplace=True)
```

```
In [14]:
# Drop rows where deck is null
df.dropna(subset=['deck'], inplace=True)
print(df.isnull().sum())
```

```
survived      0
pclass        0
sex            0
age           0
sibsp         0
parch         0
```


```
fare          0
embarked      0
class         0
who           0
adult_male    0
deck         0
embark_town   2
alive         0
alone         0
dtype: int64
```

## Step 2: Outlier Detection and Handling

- Plot boxplot and detect outliers in fare
- Cap outliers using IQR method

```
In [15]:
# Plot boxplot and detect outliers in fare
sns.boxplot(y=df['fare'])
```

<Axes: ylabel='fare'>



```
In [16]:
# Cap outliers using IQR method
df = sns.load_dataset('titanic')
Q1 = df['fare'].quantile(0.25)
Q3 = df['fare'].quantile(0.75)
IQR = Q3 - Q1
lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR
print("Values:", "\n Q1: ", Q1, "\n Q3: ", Q3, "\n IQR: ", IQR, "\n lower: ",
      lower, "\n upper: ", upper)
# df.loc[df['fare'] < lower, 'fare'] = lower
# df.loc[df['fare'] > upper, 'fare'] = upper
```

Values:  
Q1: 7.9104




```
Q3: 31.0  
IQR: 23.0896  
lower: -26.724  
upper: 65.6344
```

```
In [17]:  
# print outliers  
outliers = df[(df["fare"] < lower) | (df["fare"] > upper)]  
# print(outliers)
```

```
In [18]:  
df['fare'] = df['fare'].apply(  
    lambda X: upper if X > upper  
    else lower if X < lower  
    else X  
)
```

```
In [19]:  
sns.boxplot(y=df['fare'])
```

<Axes: ylabel='fare'>



### Step 3: Normalization and Standardization

- Apply:
- Min-Max Normalization on fare
- StandardScaler on age (use sklearn)

```
In [20]:  
# Min-Max Normalization on fare  
import from sklearn.preprocessing MinMaxScaler  
min_max_scaler = MinMaxScaler()  
# print("Print scaler fit: ",scaler.fit(df[['fare']].head(5)))  
df['Norm_fare'] = min_max_scaler.fit_transform(df[['fare']])  
print(df[['fare', 'Norm_fare']].head(5))
```

```
fare  Norm_fare
0    7.2500    0.110460
1   65.6344    1.000000
2    7.9250    0.120745
3   53.1000    0.809027
4    8.0500    0.122649
```

```
In [21]:
# StandardScaler on age (use sklearn)
import from sklearn.preprocessing StandardScaler
std_scaler = StandardScaler()
df['Std_age'] = std_scaler.fit_transform(df[['age']])
print(df[['age', 'Std_age']].head(5))
```

```
age  Std_age
0  22.0 -0.530377
1  38.0  0.571831
2  26.0 -0.254825
3  35.0  0.365167
4  35.0  0.365167
```

#### Step 4: Encoding Categorical Variables

Perform the following encodings:

- sex: Label Encoding
- embarked: One-Hot Encoding
- class: Frequency Encoding
- who: Target Encoding (target: survived)
- deck: Binary Encoding

```
In [22]:
# sex: Label Encoding
import from sklearn.preprocessing LabelEncoder
le = LabelEncoder()
df['le_sex'] = le.fit_transform(df['sex'])
print(df[['sex', 'le_sex']].head(5))
```

```
sex  le_sex
0    male      1
1  female      0
2  female      0
3  female      0
4    male      1
```

```
In [23]:
# embarked: One-Hot Encoding
dummydata = pd.get_dummies(df, columns=['embarked'])
print(dummydata[['embarked_C', 'embarked_Q', 'embarked_S']].head(5))
```

```
embarked_C  embarked_Q  embarked_S
0         False        False        True
1          True        False        False
```

2	False	False	True
3	False	False	True
4	False	False	True

```
In [24]:
# class: Frequency Encoding
freq = df['class'].value_counts()
df['freq_class'] = df['class'].map(freq)
print(df[['class', 'freq_class']].head(5))
```

```
class freq_class
0  Third      491
1  First      216
2  Third      491
3  First      216
4  Third      491
```

```
In [25]:
# who: Target Encoding (target = survived)
target_mean = df.groupby('who')['survived'].mean()
df['target_who'] = df['who'].map(target_mean)
print(df[['who', 'target_who']].head(5))
```

```
who  target_who
0   man    0.163873
1  woman    0.756458
2  woman    0.756458
3  woman    0.756458
4   man    0.163873
```

```
In [26]:
# deck: Binary Encoding
!pip install category_encoders
import category_encoders as ce

binary_encoder = ce.BinaryEncoder(cols=['deck'])
df = binary_encoder.fit_transform(df)
print(df.filter(like='deck').head(5))
```

```
Requirement already satisfied: category_encoders in
/usr/local/lib/python3.12/dist-packages (2.8.1)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (2.0.2)
Requirement already satisfied: pandas>=1.0.5 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (2.2.2)
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (1.0.1)
Requirement already satisfied: scikit-learn>=1.6.0 in
/usr/local/lib/python3.12/dist-packages (from category_encoders) (1.6.1)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.12/dist-
packages (from category_encoders) (1.16.1)
Requirement already satisfied: statsmodels>=0.9.0 in
/usr/local/lib/python3.12/dist-packages (from category_encoders) (0.14.5)
```

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.5->category\_encoders) (2.9.0.post0)  
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.5->category\_encoders) (2025.2)  
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.0.5->category\_encoders) (2025.2)  
 Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=1.6.0->category\_encoders) (1.5.1)  
 Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=1.6.0->category\_encoders) (3.6.0)  
 Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.12/dist-packages (from statsmodels>=0.9.0->category\_encoders) (25.0)  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas>=1.0.5->category\_encoders) (1.17.0)

	deck_0	deck_1	deck_2	deck_3
0	1	0	0	0
1	0	0	0	1
2	1	0	0	0
3	0	0	0	1
4	1	0	0	0

## Exercise3: Feature Selection

**Target variable: survived**

### Step 1: Pearson Correlation

- Compute correlation among: age, fare, parch, sibsp
- Drop one of any pair with correlation > 0.9

```
In [27]:
# Step 1: Pearson Correlation
num_features = ['age', 'fare', 'parch', 'sibsp']
corr_matrix = df[num_features].corr()

print("Correlation Matrix:\n", corr_matrix)

# Drop one feature if correlation > 0.9
to_drop = set()
for i in range(len(num_features)):
    for j in range(i+1, len(num_features)):
        if abs(corr_matrix.iloc[i, j]) > 0.9:
            to_drop.add(num_features[j])

print("Highly correlated features to drop:", to_drop)
```

Correlation Matrix:

	age	fare	parch	sibsp
age	1.000000	0.151920	-0.189119	-0.308247
fare	0.151920	1.000000	0.292616	0.332021
parch	-0.189119	0.292616	1.000000	0.414838
sibsp	-0.308247	0.332021	0.414838	1.000000

Highly correlated features to drop: set()

## Step 2: ANOVA (f\_classif)

- Test: sex, embarked, class, who vs survived

```
In [28]:
import sklearn.feature_selection  f_classif

cat_features = ['sex', 'embarked', 'class', 'who']
X_cat = pd.get_dummies(df[cat_features], drop_first=True)
y = df['survived']

f_values, p_values = f_classif(X_cat, y)

anova_results = pd.DataFrame({"Feature": X_cat.columns,
                              "F-value": f_values,
                              "p-value": p_values})
print("\nANOVA results:\n", anova_results)
```

ANOVA results:

	Feature	F-value	p-value
0	sex_male	372.405746	1.406055e-69
1	embarked_Q	0.011844	9.133627e-01
2	embarked_S	22.075464	3.036118e-06
3	class_Second	7.814795	5.293684e-03
4	class_Third	103.057582	5.510325e-23
5	who_man	400.037628	8.998056e-74
6	who_woman	306.864785	3.009422e-59

## Step 3: Chi-Square Test

- Apply on: sex, embarked, class vs survived

```
In [29]:
import sklearn.feature_selection  chi2

chi_features = ['sex', 'embarked', 'class']
X_chi = pd.get_dummies(df[chi_features], drop_first=True)
y = df['survived']

chi2_stats, chi2_p = chi2(X_chi, y)

chi2_results = pd.DataFrame({"Feature": X_chi.columns,
                              "Chi2": chi2_stats,
                              "p-value": chi2_p})
print("\nChi-Square results:\n", chi2_results)
```

Chi-Square results:

	Feature	Chi2	p-value
0	sex_male	92.702447	6.077838e-22
1	embarked_Q	0.010847	9.170520e-01
2	embarked_S	5.984840	1.442935e-02
3	class_Second	6.160767	1.306146e-02
4	class_Third	41.553071	1.147141e-10

## Step 4: SelectKBest

- Use SelectKBest with chi2 and f\_classif
- Select top 5 features for predicting survived

```
In [30]:
import from sklearn.feature_selection SelectKBest, f_classif, chi2
import from sklearn.impute SimpleImputer

# Separate numerical & categorical features
num_features = ['age', 'fare', 'parch', 'sibsp']
X_num = df[num_features]
X_cat = pd.get_dummies(df[['sex', 'embarked', 'class', 'who']],
drop_first=True).astype(int)

# Handle missing values
imputer_num = SimpleImputer(strategy='mean')

X_num = pd.DataFrame(imputer_num.fit_transform(X_num), columns=num_features)
# (no need to impute X_cat now, as dummies are 0/1 integers without NaNs)

# Final dataset
X = pd.concat([X_num, X_cat], axis=1)
y = df['survived']

# --- SelectKBest with f_classif ---
selector_f = SelectKBest(score_func=f_classif, k=5)
X_new_f = selector_f.fit_transform(X, y)
print("\nTop 5 features by ANOVA f_classif:\n",
X.columns[selector_f.get_support()].tolist())

# --- SelectKBest with chi2 ---
selector_chi = SelectKBest(score_func=chi2, k=5)
X_new_chi = selector_chi.fit_transform(abs(X), y) # chi2 needs non-negative
values
print("\nTop 5 features by Chi-Square:\n",
X.columns[selector_chi.get_support()].tolist())
```

Top 5 features by ANOVA f\_classif:

['fare', 'sex\_male', 'class\_Third', 'who\_man', 'who\_woman']

Top 5 features by Chi-Square:

['fare', 'sex\_male', 'class\_Third', 'who\_man', 'who\_woman']

```
In [31]:
import from sklearn.feature_selection SelectKBest, f_classif, chi2
import from sklearn.impute SimpleImputer
import pandas as pd

# Separate numerical & categorical features
num_features = ['age', 'fare', 'parch', 'sibsp']
X_num = df[num_features]
X_cat = pd.get_dummies(df[['sex', 'embarked', 'class', 'who']],
drop_first=True).astype(int)

# Handle missing values
imputer_num = SimpleImputer(strategy='mean')
X_num = pd.DataFrame(imputer_num.fit_transform(X_num), columns=num_features)
```

```

# Final dataset
X = pd.concat([X_num, X_cat], axis=1)
y = df['survived']

# --- Select top 5 with f_classif ---
selector_f = SelectKBest(score_func=f_classif, k=5)
selector_f.fit(X, y)
f_results = pd.DataFrame({
    "Feature": X.columns,
    "Score (F)": selector_f.scores_,
    "p-value": selector_f.pvalues_
}).sort_values(by="Score (F)", ascending=False)

print("\nTop 5 features by ANOVA f_classif:\n", f_results.head(5))

# --- Select top 5 with chi2 ---
selector_chi = SelectKBest(score_func=chi2, k=5)
selector_chi.fit(abs(X), y) # chi2 needs non-negative values
chi_results = pd.DataFrame({
    "Feature": X.columns,
    "Score (Chi2)": selector_chi.scores_,
    "p-value": selector_chi.pvalues_
}).sort_values(by="Score (Chi2)", ascending=False)

print("\nTop 5 features by Chi-Square:\n", chi_results.head(5))

```

Top 5 features by ANOVA f\_classif:

	Feature	Score (F)	p-value
9	who_man	400.037563	8.998259e-74
4	sex_male	372.405724	1.406066e-69
10	who_woman	306.864782	3.009426e-59
8	class_Third	103.057599	5.510281e-23
1	fare	99.614870	2.622660e-22

Top 5 features by Chi-Square:

	Feature	Score (Chi2)	p-value
1	fare	1564.437823	0.000000e+00
10	who_woman	159.095048	1.783910e-36
9	who_man	109.859713	1.051769e-25
4	sex_male	92.702447	6.077838e-22
8	class_Third	41.553071	1.147141e-10