

## Lab 5 Supervised Learning– Regression

```
In [1]:
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn.datasets fetch_california_housing
```

```
In [2]:
# dataset
housing = fetch_california_housing(as_frame=True)
df = housing.frame
```

```
In [3]:
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   MedInc          20640 non-null  float64
1   HouseAge        20640 non-null  float64
2   AveRooms        20640 non-null  float64
3   AveBedrms       20640 non-null  float64
4   Population      20640 non-null  float64
5   AveOccup        20640 non-null  float64
6   Latitude        20640 non-null  float64
7   Longitude       20640 non-null  float64
8   MedHouseVal     20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```

```
In [4]:
# print(df.head(5))
```

### Exercises

**Apply the following regression models to predict charges:**

- Simple Linear Regression (using MedInc only)
- Multiple Linear Regression (all features)
- Polynomial Regression (degree 2 and 3, using both single and multiple

features)

- Ridge Regression (with tuning alpha)
- Lasso Regression (with tuning alpha)
- Decision Tree Regressor

**For each model:**

- Train the model
- Predict on test data
- Evaluate and record:
  - MAE
  - MSE
  - RMSE
  - $R^2$  Score
- Create scatter plot of actual vs predicted value

```
In [5]:
import from sklearn.model_selection train_test_split
import from sklearn.metrics mean_absolute_error, mean_squared_error, r2_score
import from sklearn.linear_model LinearRegression
```

```
In [6]:
# features and target
X = df.drop('MedHouseVal', axis =1)
y = df['MedHouseVal']
```

```
In [7]:
# Train test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=12)
```

```
In [8]:
# evaluate functions
def evaluate(y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    print(f'MAE: {mae:.2f}')
    print(f'MSE: {mse:.2f}')
    print(f'RMSE: {rmse:.2f}')
    print(f'R2 Score: {r2:.2f}')
```

```
In [9]:
# visualise funtion
def visualize(y_test, y_pred):
    plt.scatter(y_test[y_pred > y_test], y_pred[y_pred > y_test], color='red',
label='Predicted > Actual')
    plt.scatter(y_test[y_pred < y_test], y_pred[y_pred < y_test], color='blue',
label='Predicted < Actual')

    # Add a line representing the linear regression fit
    m, b = np.polyfit(y_test, y_pred, 1)
    plt.plot(y_test, m*y_test + b, color='green', label='Linear Fit')
```

```
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs Predicted')
plt.legend()
plt.show()
```

### Simple Linear Regression (using MedInc only)

```
In [10]:
# train model - simple linear regression
X_train_inc = X_train[['MedInc']]
X_test_inc = X_test[['MedInc']]

model = LinearRegression()
model.fit(X_train_inc, y_train)
```

LinearRegression()

```
In [11]:
prediction = model.predict(X_test_inc)
```

```
In [12]:
print("Simple Linear Regression:")
evaluate(y_test, prediction)
visualize(y_test, prediction)
```


Simple Linear Regression:

MAE: 0.63

MSE: 0.70

RMSE: 0.84

R2 Score: 0.49



### Multiple Linear Regression (all features)


```
In [13]:  
# train model - multiple linear regression  
X_train_mult = X_train  
X_test_mult = X_test  
  
model = LinearRegression()  
model.fit(X_train_mult, y_train)
```

LinearRegression()

```
In [14]:  
prediction = model.predict(X_test)
```

```
In [15]:  
print("Multiple Linear Regression:")  
evaluate(y_test, prediction)  
visualize(y_test, prediction)
```

Multiple Linear Regression:  
MAE: 0.53  
MSE: 0.53  
RMSE: 0.73  
R2 Score: 0.62



### Polynomial Regression (degree 2 and 3, using both single and multiple features)

```
In [16]:  
import from sklearn.preprocessing PolynomialFeatures  
import from sklearn.pipeline make_pipeline
```

```
In [17]:  
# degree 2 simple regression  
model = make_pipeline(PolynomialFeatures(2), LinearRegression())  
model.fit(X_train_inc, y_train)  
prediction = model.predict(X_test_inc)
```

```
print("Polynomial Regression (degree 2):")
evaluate(y_test, prediction)
# visualize(y_test, prediction)
```

Polynomial Regression (degree 2):  
MAE: 0.62  
MSE: 0.69  
RMSE: 0.83  
R2 Score: 0.49

```
In [18]:
# degree 2 simple regression
model = make_pipeline(PolynomialFeatures(2), LinearRegression())
model.fit(X_train, y_train)
prediction = model.predict(X_test)
print("Polynomial Regression (degree 2):")
evaluate(y_test, prediction)
# visualize(y_test, prediction)
```

Polynomial Regression (degree 2):  
MAE: 0.46  
MSE: 0.44  
RMSE: 0.66  
R2 Score: 0.68

```
In [19]:
# degree 3 simple regression
model = make_pipeline(PolynomialFeatures(3), LinearRegression())
model.fit(X_train_inc, y_train)
prediction = model.predict(X_test_inc)
print("Polynomial Regression (degree 3):")
evaluate(y_test, prediction)
# visualize(y_test, prediction)
```

Polynomial Regression (degree 3):  
MAE: 0.62  
MSE: 0.69  
RMSE: 0.83  
R2 Score: 0.50

```
In [20]:
# degree 3 multi regression
model = make_pipeline(PolynomialFeatures(3), LinearRegression())
model.fit(X_train, y_train)
prediction = model.predict(X_test)
print("Polynomial Regression (degree 3):")
evaluate(y_test, prediction)
# visualize(y_test, prediction)
```

Polynomial Regression (degree 3):  
MAE: 0.44  
MSE: 0.44  
RMSE: 0.66  
R2 Score: 0.68

```

In [21]:
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Degree 2 Simple Regression
model_deg2_simple = make_pipeline(PolynomialFeatures(2), LinearRegression())
model_deg2_simple.fit(X_train_inc, y_train)
prediction_deg2_simple = model_deg2_simple.predict(X_test_inc)
axes[0, 0].scatter(y_test[prediction_deg2_simple > y_test],
prediction_deg2_simple[prediction_deg2_simple > y_test], color='red',
label='Predicted > Actual')
axes[0, 0].scatter(y_test[prediction_deg2_simple < y_test],
prediction_deg2_simple[prediction_deg2_simple < y_test], color='blue',
label='Predicted < Actual')
m, b = np.polyfit(y_test, prediction_deg2_simple, 1)
axes[0, 0].plot(y_test, m*y_test + b, color='green', label='Linear Fit')
axes[0, 0].set_title('Polynomial Regression (Degree 2, Simple)')
axes[0, 0].set_xlabel('Actual')
axes[0, 0].set_ylabel('Predicted')
axes[0, 0].legend()

# Degree 2 Multiple Regression
model_deg2_mult = make_pipeline(PolynomialFeatures(2), LinearRegression())
model_deg2_mult.fit(X_train_mult, y_train)
prediction_deg2_mult = model_deg2_mult.predict(X_test_mult)
axes[0, 1].scatter(y_test[prediction_deg2_mult > y_test],
prediction_deg2_mult[prediction_deg2_mult > y_test], color='red',
label='Predicted > Actual')
axes[0, 1].scatter(y_test[prediction_deg2_mult < y_test],
prediction_deg2_mult[prediction_deg2_mult < y_test], color='blue',
label='Predicted < Actual')
m, b = np.polyfit(y_test, prediction_deg2_mult, 1)
axes[0, 1].plot(y_test, m*y_test + b, color='green', label='Linear Fit')
axes[0, 1].set_title('Polynomial Regression (Degree 2, Multiple)')
axes[0, 1].set_xlabel('Actual')
axes[0, 1].set_ylabel('Predicted')
axes[0, 1].legend()

# Degree 3 Simple Regression
model_deg3_simple = make_pipeline(PolynomialFeatures(3), LinearRegression())
model_deg3_simple.fit(X_train_inc, y_train)
prediction_deg3_simple = model_deg3_simple.predict(X_test_inc)
axes[1, 0].scatter(y_test[prediction_deg3_simple > y_test],
prediction_deg3_simple[prediction_deg3_simple > y_test], color='red',
label='Predicted > Actual')
axes[1, 0].scatter(y_test[prediction_deg3_simple < y_test],
prediction_deg3_simple[prediction_deg3_simple < y_test], color='blue',
label='Predicted < Actual')
m, b = np.polyfit(y_test, prediction_deg3_simple, 1)
axes[1, 0].plot(y_test, m*y_test + b, color='green', label='Linear Fit')
axes[1, 0].set_title('Polynomial Regression (Degree 3, Simple)')
axes[1, 0].set_xlabel('Actual')
axes[1, 0].set_ylabel('Predicted')
axes[1, 0].legend()

# Degree 3 Multiple Regression
model_deg3_mult = make_pipeline(PolynomialFeatures(3), LinearRegression())
model_deg3_mult.fit(X_train_mult, y_train)
prediction_deg3_mult = model_deg3_mult.predict(X_test_mult)

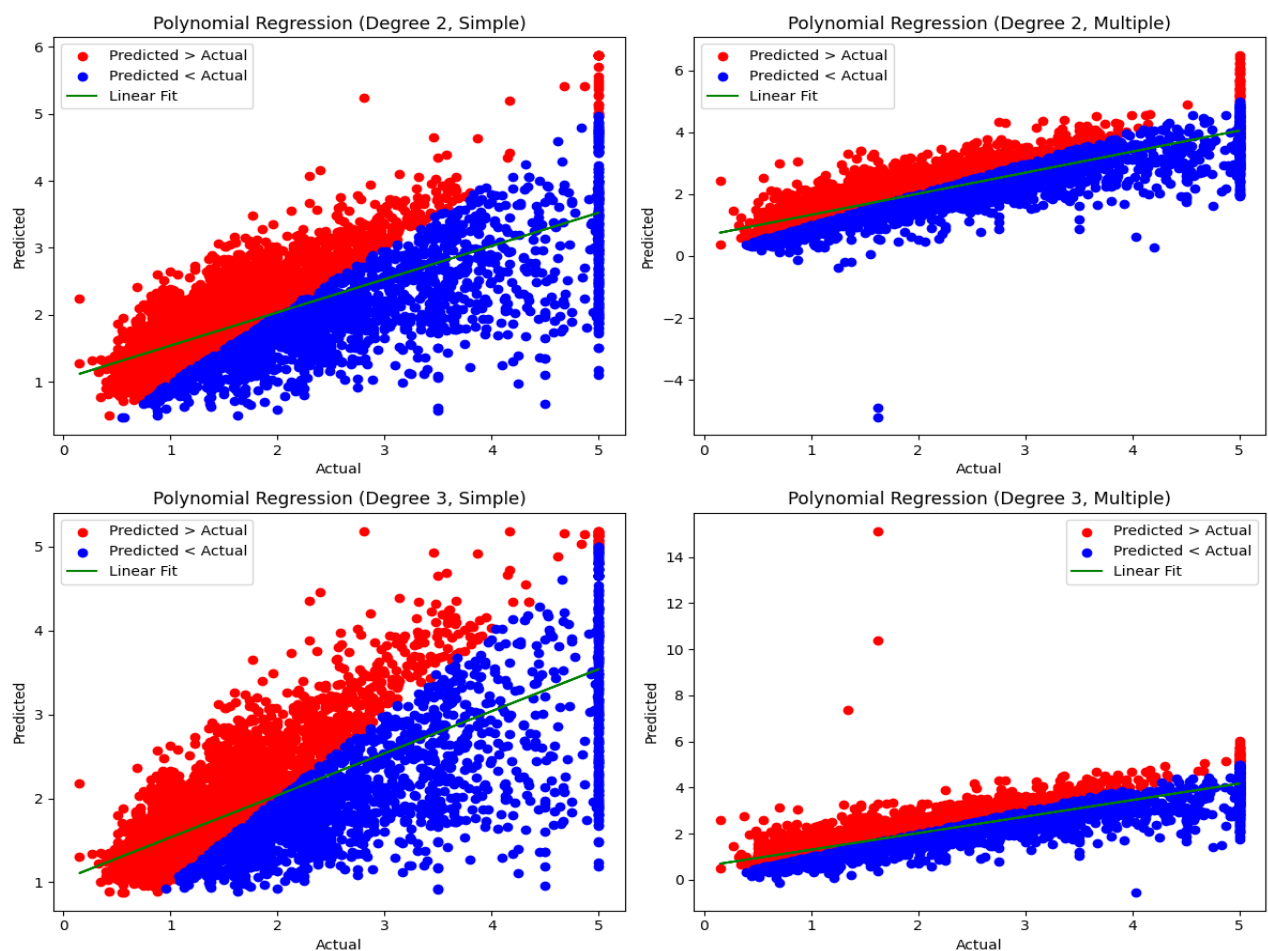
```

```

axes[1, 1].scatter(y_test[prediction_deg3_mult > y_test],
prediction_deg3_mult[prediction_deg3_mult > y_test], color='red',
label='Predicted > Actual')
axes[1, 1].scatter(y_test[prediction_deg3_mult < y_test],
prediction_deg3_mult[prediction_deg3_mult < y_test], color='blue',
label='Predicted < Actual')
m, b = np.polyfit(y_test, prediction_deg3_mult, 1)
axes[1, 1].plot(y_test, m*y_test + b, color='green', label='Linear Fit')
axes[1, 1].set_title('Polynomial Regression (Degree 3, Multiple)')
axes[1, 1].set_xlabel('Actual')
axes[1, 1].set_ylabel('Predicted')
axes[1, 1].legend()

plt.tight_layout()
plt.show()

```



### Ridge Regression (with tuning alpha)

```

In [22]:
import sklearn.linear_model as Ridge

model = Ridge()
model.fit(X_train, y_train)

prediction = model.predict(X_test)

```

```
print("Ridge Regression:")  
evaluate(y_test, prediction)  
visualize(y_test, prediction)
```


Ridge Regression:

MAE: 0.53

MSE: 0.53

RMSE: 0.73

R2 Score: 0.62



### Lasso Regression (with tuning alpha)

```
In [23]:  
import sklearn.linear_model  Lasso  
  
model = Lasso()  
model.fit(X_train, y_train)  
  
prediction = model.predict(X_test)  
  
print("Lasso Regression:")  
evaluate(y_test, prediction)  
visualize(y_test, prediction)
```

Lasso Regression:


MAE: 0.78

MSE: 0.98

RMSE: 0.99

R2 Score: 0.29





### Decision Tree Regressor

```
In [24]:
import sklearn.tree DecisionTreeRegressor

model = DecisionTreeRegressor()
model.fit(X_train, y_train)

prediction = model.predict(X_test)

print("Decision Tree Regressor:")
evaluate(y_test, prediction)
visualize(y_test, prediction)
```

Decision Tree Regressor:

MAE: 0.44

MSE: 0.47

RMSE: 0.68

R2 Score: 0.66

